

In [1]:

```
#Examining Factors Responsible for Heart Attacks
```

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="whitegrid")
```

In [3]:

```
#Section 1 - Loading the data and exploratory data analysis
```

In [4]:

```
#Load the data
```

In [7]:

```
data = pd.read_csv("heart_Data.csv")
data.head()
```

Out[7]:

| Unnamed: 0 | age | sex | chest_pain_type | resting_bp | cholesterol | fasting_blood_sugar | restecg | max_hr | exang | oldpeak | slope |
|------------|-----|-----|-----------------|------------|-------------|---------------------|---------|--------|-------|---------|-------|
| 0          | 0   | 63  | 1               | 3          | 145         | 233                 | 1       | 0      | 150   | 0       | 2.3   |
| 1          | 1   | 37  | 1               | 2          | 130         | 250                 | 0       | 1      | 187   | 0       | 3.5   |
| 2          | 2   | 41  | 0               | 1          | 130         | 204                 | 0       | 0      | 172   | 0       | 1.4   |
| 3          | 3   | 56  | 1               | 1          | 120         | 236                 | 0       | 1      | 178   | 0       | 0.8   |
| 4          | 4   | 57  | 0               | 0          | 120         | 354                 | 0       | 1      | 163   | 1       | 0.6   |

In [12]:

```
#data.to_csv("heart_Data")
```

**Performing EDA: 4. Identify the data variables which might be categorical in nature. Describe and explore these variables using appropriate tools. For example: count plot.**

In [9]:

```
data.shape
```


Out[9]:

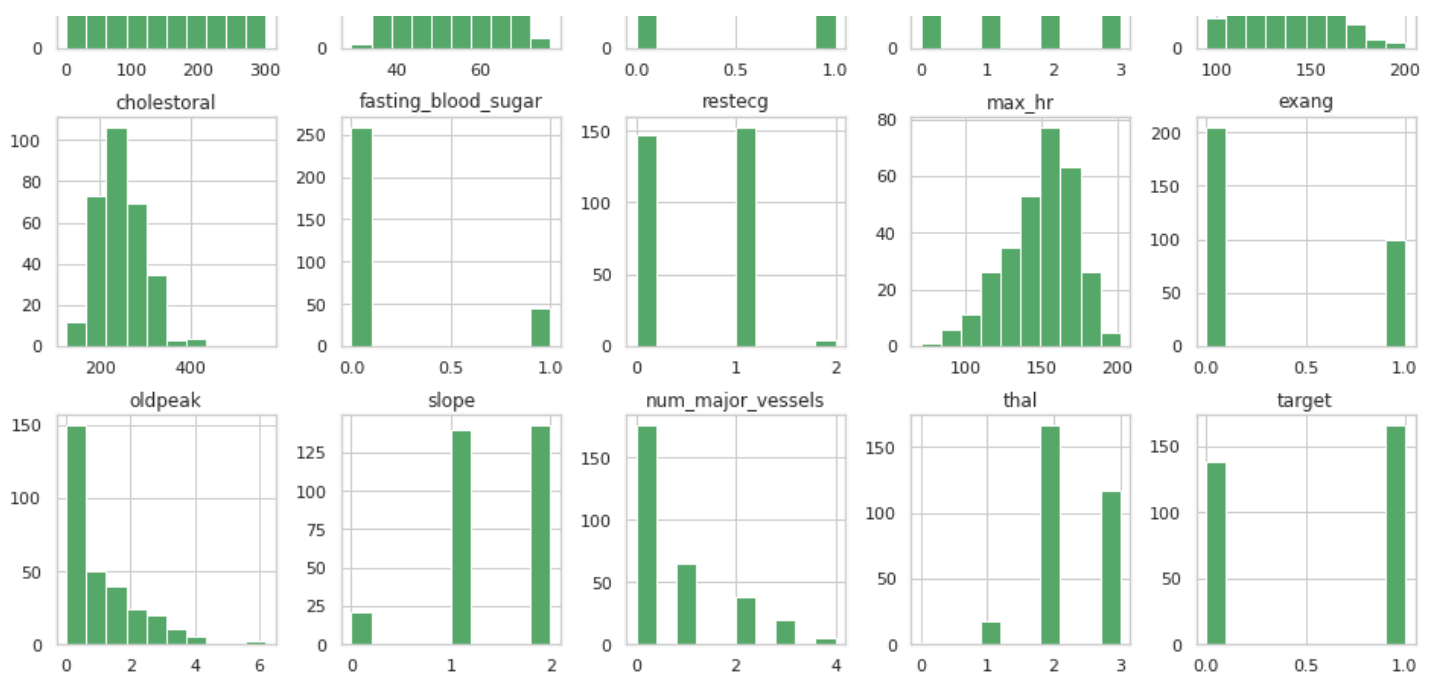
```
(303, 15)
```

In [10]:

```
data.hist(layout = (3,5), figsize=(16,10), color = 'g')
print('Data Distribution')
```

Data Distribution






In [11]:


```
print('This looks like a fairly balanced dataset, as distribution of majority and minority class is around 55:45')
sns.countplot(x="target", data=data, palette="mako_r")
plt.show()
```

This looks like a fairly balanced dataset, as distribution of majority and minority class is around 55:45



In [13]:

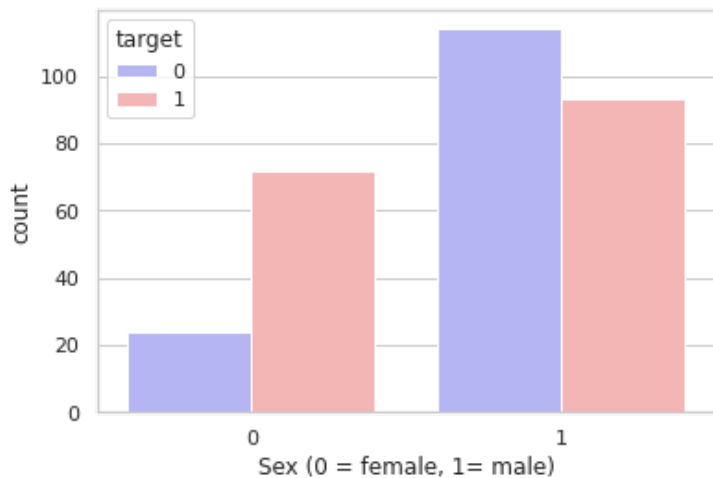
```
sns.countplot(x='sex', data=data, palette="bwr")
plt.xlabel("Sex (0 = female, 1= male)")
plt.show()
```



In [14]:


```
print('Analysing distribution of target and sex (0-female 1-male)')
sns.countplot(x = data['sex'], hue = data['target'], palette='bwr')
plt.xlabel("Sex (0 = female, 1= male)")
plt.show()
```

Analysing distribution of target and sex (0-female 1-male)




In [15]:


```
pd.crosstab(data.age, data.target).plot(kind="bar", figsize=(20,6), color = ['g', 'r'])
plt.title('Heart Disease Distribution by Patient Age')
plt.xlabel('Age')
plt.ylabel('Counts')
plt.show()
```



In [16]:

```
plt.bar(data.age[data.target==1], data.max_hr[(data.target==1)], color="red")
plt.bar(data.age[data.target==0], data.max_hr[(data.target==0)], color="grey")
plt.legend(["Diseased", "Not Diseased"])
plt.xlabel("Age")
plt.ylabel("Maximum Heart Rate")
plt.show()
```





In [17]:

```
## Section 2 - Data Pre-processing
```

In [18]:

```
#Null or missing value check
```

In [19]:

```
data.isnull().sum()
```

Out[19]:

```
Unnamed: 0      0
age             0
sex            0
chest_pain_type 0
resting_bp      0
cholestorol     0
fasting_blood_sugar 0
restecg        0
max_hr         0
exang          0
oldpeak        0
slope          0
num_major_vessels 0
thal           0
target         0
dtype: int64
```

**We do not see any missing values for this dataset.**

In [20]:

```
#Duplicate inspection
```

In [21]:

```
data.duplicated().any()
```

Out[21]:

```
False
```

**Looks like the dataset has some duplicates. Let's remove the duplicates**

In [22]:

```
data.drop_duplicates(subset=None, inplace=True)
data.duplicated().any()
```

Out[22]:

```
False
```

In [23]:

```
data.shape
```

Out[23]:

```
(303, 15)
```

**So, we can see that there was one duplicate row**

In [24]:

```
#One Hot Encoding
```

In [25]:

```
def encode_features(df, features):  
    '''  
    Method for one-hot encoding all selected categorical fields  
    '''  
    for f in features:  
        if f in df.columns:  
            encoded = pd.get_dummies(df[f])  
            encoded = encoded.add_prefix(f + '_')  
            df = pd.concat([df, encoded], axis=1)  
        else:  
            print('Feature not found')  
            return df  
  
    df.drop(columns=features, inplace = True)  
  
    return df
```

In [26]:

```
features_to_encode = ['thal', 'slope', 'chest_pain_type', 'restecg']  
encoded = encode_features(data, features_to_encode)  
data = encoded.copy()  
print(data.shape)
```

Feature not found  
(303, 15)

In [27]:

```
data.columns
```

Out[27]:

```
Index(['Unnamed: 0', 'age', 'sex', 'resting_bp', 'cholestorl',  
      'fasting_blood_sugar', 'max_hr', 'exang', 'oldpeak',  
      'num_major_vessels', 'target', 'thal_0', 'thal_1', 'thal_2', 'thal_3'],  
      dtype='object')
```

In [28]:

```
#Outlier Inspection
```

In [30]:

```
data.describe()
```

Out[30]:

|       | Unnamed:<br>0 | age        | sex        | resting_bp | cholestorl | fasting_blood_sugar | max_hr     | exang      | oldpeak    |
|-------|---------------|------------|------------|------------|------------|---------------------|------------|------------|------------|
| count | 303.000000    | 303.000000 | 303.000000 | 303.000000 | 303.000000 | 303.000000          | 303.000000 | 303.000000 | 303.000000 |
| mean  | 151.000000    | 54.366337  | 0.683168   | 131.623762 | 246.264026 | 0.148515            | 149.646865 | 0.326733   | 1.039604   |
| std   | 87.612784     | 9.082101   | 0.466011   | 17.538143  | 51.830751  | 0.356198            | 22.905161  | 0.469794   | 1.161075   |
| min   | 0.000000      | 29.000000  | 0.000000   | 94.000000  | 126.000000 | 0.000000            | 71.000000  | 0.000000   | 0.000000   |
| 25%   | 75.500000     | 47.500000  | 0.000000   | 120.000000 | 211.000000 | 0.000000            | 133.500000 | 0.000000   | 0.000000   |
| 50%   | 151.000000    | 55.000000  | 1.000000   | 130.000000 | 240.000000 | 0.000000            | 153.000000 | 0.000000   | 0.800000   |
| 75%   | 226.500000    | 61.000000  | 1.000000   | 140.000000 | 274.500000 | 0.000000            | 166.000000 | 1.000000   | 1.600000   |
| max   | 302.000000    | 77.000000  | 1.000000   | 200.000000 | 564.000000 | 1.000000            | 202.000000 | 1.000000   | 6.200000   |

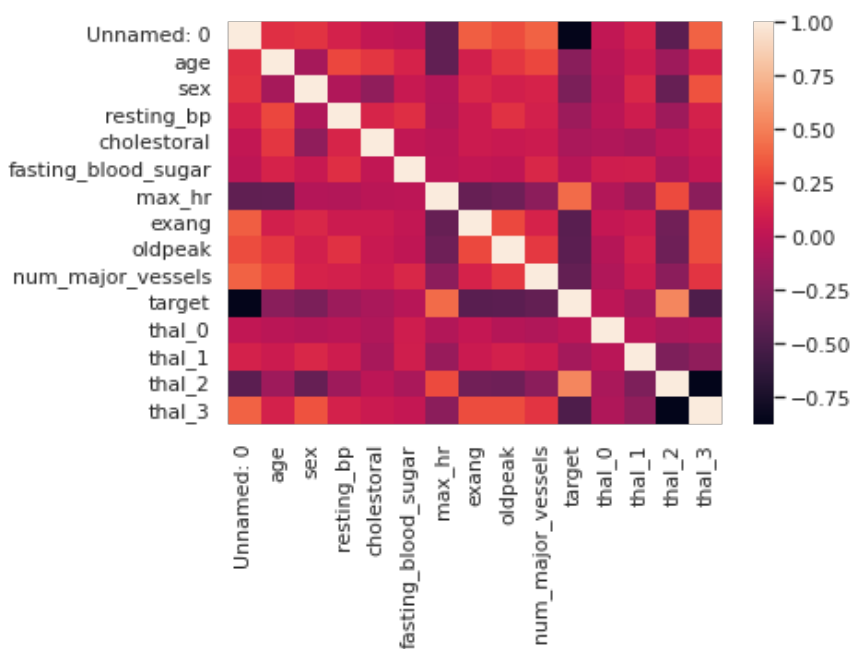
In [31]:

```
# Data Correlation
```

In [32]:

```
print(data.corr()['target'])
sns.heatmap(data.corr())
plt.show()
```

```
Unnamed: 0      -0.862585
age             -0.225439
sex             -0.280937
resting_bp      -0.144931
cholestorl      -0.085239
fasting_blood_sugar -0.028046
max_hr          0.421741
exang           -0.436757
oldpeak         -0.430696
num_major_vessels -0.391724
target          1.000000
thal_0          -0.007293
thal_1          -0.106589
thal_2          0.527334
thal_3          -0.486112
Name: target, dtype: float64
```



In [34]:

```
#Section-3 : Feature Engineering
```

In [35]:

```
feature_engg_data = data.copy()
outlier_data = data.copy()
target_index = data.columns.get_loc("target") # Use this for separating the target labels from data
factor = 3
# Include this only for columns with suspected outliers
# Using a factor of 3, following Nelson's rule 1 to remove outliers - https://en.wikipedia.org/wiki/Nelson_rules
# Only for non-categorical fields
columns_to_include = ['age', 'resting_bp', 'cholestorl', 'max_hr', 'oldpeak', 'num_major_vessels']
for column in columns_to_include:
    upper_lim = feature_engg_data[column].mean () + feature_engg_data[column].std () * factor
    lower_lim = feature_engg_data[column].mean () - feature_engg_data[column].std () * factor
    feature_engg_data = feature_engg_data[(feature_engg_data[column] < upper_lim) & (feature_engg_data[column] > lower_lim)]

outlier_data= pd.concat([outlier_data, feature_engg_data]).drop_duplicates(keep=False)
```

In [36]:

```
print(feature_engg_data.shape)
print(outlier_data.shape)
```

```
(289, 15)
(14, 15)
```

In [37]:

```
#Data Normalization
```

In [38]:

```
from sklearn import preprocessing

def normalize_data(df):
    val = df.values
    min_max_normalizer = preprocessing.MinMaxScaler()
    norm_val = min_max_normalizer.fit_transform(val)
    df2 = pd.DataFrame(norm_val)

    return df2

norm_feature_engg_data = normalize_data(feature_engg_data)
norm_outlier_data = normalize_data(outlier_data)
```

In [39]:

```
norm_feature_engg_data = normalize_data(feature_engg_data)
norm_outlier_data = normalize_data(outlier_data)
```

**#Data Splits Splitting Feature Engineered Data into train-valid-test dataset in 70:20:10 Ratio, the choice of selecting this splitting ratio is to ensure we have sufficient training data, sufficient validation data for mainly hyper parameter tuning and sufficient testing data to ensure model generalization.**

In [40]:

```
from sklearn.model_selection import train_test_split

input_data = norm_feature_engg_data.drop([target_index],axis='columns')
targets =norm_feature_engg_data.filter([target_index],axis='columns')

x, x_test, y, y_test = train_test_split(input_data,targets,test_size=0.1,train_size=0.9,
random_state=5)
x_train, x_valid, y_train, y_valid = train_test_split(x,y,test_size = 0.22,train_size =0
.78, random_state=5)
```

In [41]:

```
#Section 4 - Building the Model
```

In [42]:

```
#Model Evaluation Metrics
```

In [43]:

```
from sklearn.metrics import accuracy_score,roc_auc_score,confusion_matrix
import math
import seaborn as sns
```

In [44]:

```
def evaluatation_metrics(y_true, y_pred,model):

    accuracy = accuracy_score(y_true, y_pred)
    roc_auc = roc_auc_score(y_true, y_pred, average='weighted')
    cm = confusion_matrix(y_true, y_pred)
```

```

print("Accuracy of",model,": {:.2f}".format(accuracy))
print("ROC AUC Score of", model,": {:.2f}".format(roc_auc))
print("Confusion Matrix of", model,": \n")

plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues')
;

plt.ylabel('Actual label');
plt.xlabel('Predicted label');
title = 'AUC-ROC Score: {:.2f}'.format(roc_auc)
plt.title(title)
plt.show()

```

In [45]:

```

import time

def ml_model(classifier, classifier_name, **kwargs):
    """
    Generic method to train the selected classification algorithm on train, validation and test dataset.
    """
    # Fit model
    if kwargs['x_train'] is not None:
        model = classifier.fit(kwargs['x_train'], kwargs['y_train'])
        y_pred_train= model.predict(kwargs['x_train'])
        print('*****')
        print('Training Set Performance:')
        print('*****')
        evaluation_metrics(kwargs['y_train'], y_pred_train, classifier_name)

    if kwargs['x_valid'] is not None:
        y_pred_valid = model.predict(kwargs['x_valid'])
        print('*****')
        print('Validation Set Performance:')
        print('*****')
        evaluation_metrics(kwargs['y_valid'], y_pred_valid, classifier_name)

    if kwargs['x_test'] is not None:
        start = time.time()
        y_pred_test= classifier.predict(kwargs['x_test'])
        end = time.time()
        print('*****')
        print('Test Set Performance:')
        print('*****')
        print('Model Time Complexity on Test Data: {:.3f} milli seconds'.format((end - start) * 1000))
        evaluation_metrics(kwargs['y_test'], y_pred_test, classifier_name)

```

In [46]:

```

from sklearn.model_selection import learning_curve
from sklearn.model_selection import ShuffleSplit

def plot_learning_curves(train_sizes, train_scores_mean, train_scores_std, test_scores_mean, test_scores_std, model_name):
    """
    Method to generate learning curves for using training and cross validation scores
    """
    plt.title(model_name)

    plt.xlabel("Training examples")
    plt.ylabel("Score")

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1,
                     color="g")

```



```

plt.plot(train_sizes, train_scores_mean, 'o-', color="r", label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g", label="Cross-validation score")

plt.legend(loc="best")
plt.show()

def plot_model_scalability_curves(train_sizes, training_time_mean, training_time_std, model_name):
    """
    Method to generate scalability curve to see the model complexity
    """
    plt.plot(train_sizes, training_time_mean, 'o-', color = 'purple')
    plt.fill_between(train_sizes, training_time_mean - training_time_std,
                     training_time_mean + training_time_std, alpha=0.1, color = 'purple')
    plt.xlabel("Training examples")
    plt.ylabel("Training time")
    plt.title("Scalability of " + model_name)
    plt.show()

def plot_model_performance_curves(training_time_mean, test_scores_mean, test_scores_std, model_name):
    """
    Method to generate performance curves to see if increase model complexity would improve score or not
    """
    plt.plot(training_time_mean, test_scores_mean, 'o-')
    plt.fill_between(training_time_mean, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1)
    plt.xlabel("Training Time")
    plt.ylabel("Score")
    plt.title("Performance of " + model_name)
    plt.show()

def generate_learning_curves(model, model_name, X, y, xlim = None, ylim=None, epochs =None, figsize = (20,5)):
    """
    Generic method to generate Learning Curves, Scalability curves and Performance curves
    Referred - https://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_learning\_curve.html#sphx-glr-auto-examples-model-selection-plot-learning-curve-py
    """
    cross_valid = ShuffleSplit(n_splits=100, test_size=0.2, random_state=0)

    train_sizes=np.linspace(.1, 1.0, 5)
    train_scores, test_scores, training_time, _ = learning_curve(model, X, y, cv=cross_valid,
                                                                    n_jobs=epoch
s, train_sizes=train_sizes,
                                                                    return_times
=True)

    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)

    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

    training_time_mean = np.mean(training_time, axis=1)
    training_time_std = np.std(training_time, axis=1)

    # Plot learning curve
    plot_learning_curves(train_sizes, train_scores_mean, train_scores_std, test_scores_mean, test_scores_std, model_name)

    # Plot scalability curve
    plot_model_scalability_curves(train_sizes, training_time_mean, training_time_std, model_name)

    # Plot model performance score
    plot_model_performance_curves(training_time_mean, test_scores_mean, test_scores_std, model_name)

```

In [47]:

```
#Classification Algorithms
```

In [48]:

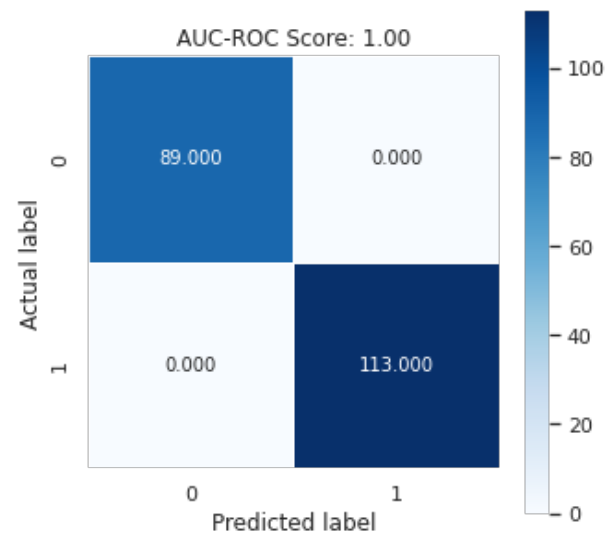
```
#Baseline - Decision Tree
```

In [49]:


```
from sklearn.tree import DecisionTreeClassifier, plot_tree

DTC = DecisionTreeClassifier(criterion='entropy', random_state=3) # Baseline model witho
ut any form of pruning
ml_model(DTC, 'Decision Tree', x_train = x_train, y_train = y_train, x_valid = x_valid,
y_valid = y_valid, x_test = None)
```

```
*****
Training Set Performance:
*****
Accuracy of Decision Tree : 1.00
ROC AUC Score of Decision Tree : 1.00
Confusion Matrix of Decision Tree :
```



```
*****
Validation Set Performance:
*****
Accuracy of Decision Tree : 1.00
ROC AUC Score of Decision Tree : 1.00
Confusion Matrix of Decision Tree :
```



In [50]:

```
# Visualizing Decision Trees
fig, axes = plt.subplots(figsize = (12,10), dpi=500)
plot_tree(DTC, filled = True)
plt.show()
```


$x[0] \leq 0.543$   
entropy = 0.99  
samples = 202  
value = [89, 113]

entropy = 0.0  
samples = 113  
value = [0, 113]

entropy = 0.0  
samples = 89  
value = [89, 0]

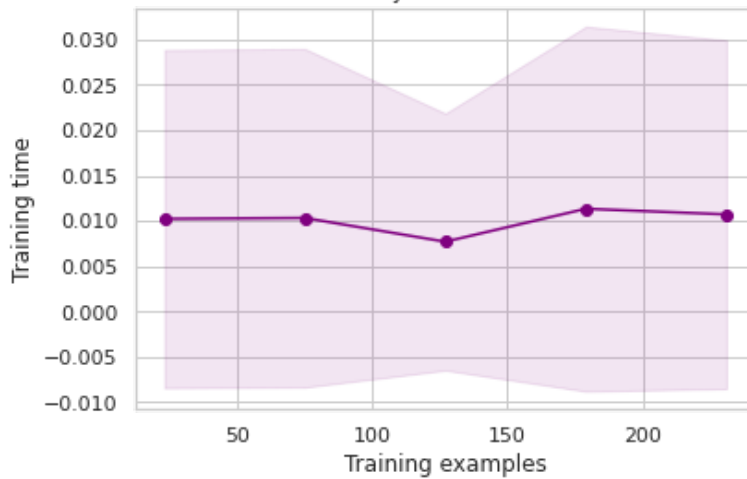
In [52]:

```
# Generation model curves on training and cross validation data
generate_learning_curves(
    model = DTC,
    model_name = "Decision Trees",
    X = input_data,
    y = targets,
    ylim=(0.7, 1.01),
    epochs=5)
```




## Training examples

Scalability of Decision Trees



Performance of Decision Trees



In [53]:

```
#Hyper-parameter tuning
```

In [54]:

```
# We will apply grid search to find the best possible set of hyper parameters
def DTGridSearch(X,y,param_lim):
    """
    Decision Tree Grid Search to find the initial best guess of hyper-parameters
    """
    param_grid = {
        'min_samples_leaf':np.linspace(param_lim[0],param_lim[1],20).astype(
'int'),
        'max_depth':np.arange(1,param_lim[2]),
        'criterion' : ['entropy' , 'gini']
    }

    best_DT = GridSearchCV(estimator = DecisionTreeClassifier(random_state=3), param_grid=param_grid, cv=10)
    best_DT.fit(X, y)

    print("Best Decision Tree Hyper-Parameters are:")
    print(best_DT.best_params_)

    return best_DT.best_params_['min_samples_leaf'], best_DT.best_params_['max_depth'],
best_DT.best_params_['criterion']
```

In [56]:

```
min_samples_leaf_lim = int(0.005 * len(x_train)) # 0.5% of length of training size
max_samples_leaf_lim = int(0.1 * len(x_train)) # 10% of length of training size
max_depth = 10
```

```
best_min_sample_leaf, best_max_depth, best_criterion = DTGridSearch( x_train,
                                                                    y_train,
                                                                    (min_samples_leaf_
lim, max_samples_leaf_lim, max_depth)
                                                                    )
```

```
-----
NameError                                Traceback (most recent call last)
/tmp/ipykernel_162/1327311137.py in <cell line: 5>()
      3 max_depth = 10
      4
----> 5 best_min_sample_leaf, best_max_depth, best_criterion = DTGridSearch( x_train,
      6                                                                    y_train,
      7                                                                    (min_sampl
es_leaf_lim, max_samples_leaf_lim, max_depth)

/tmp/ipykernel_162/3665654205.py in DTGridSearch(X, y, param_lim)
     10     }
     11
--> 12     best_DT = GridSearchCV(estimator = DecisionTreeClassifier(random_state=3), pa
ram_grid=param_grid, cv=10)
     13     best_DT.fit(X, y)
     14

NameError: name 'GridSearchCV' is not defined
```

**Best Decision Tree Hyper-Parameters are: {'criterion': 'entropy', 'max\_depth': 1, 'min\_samples\_leaf': 1}**

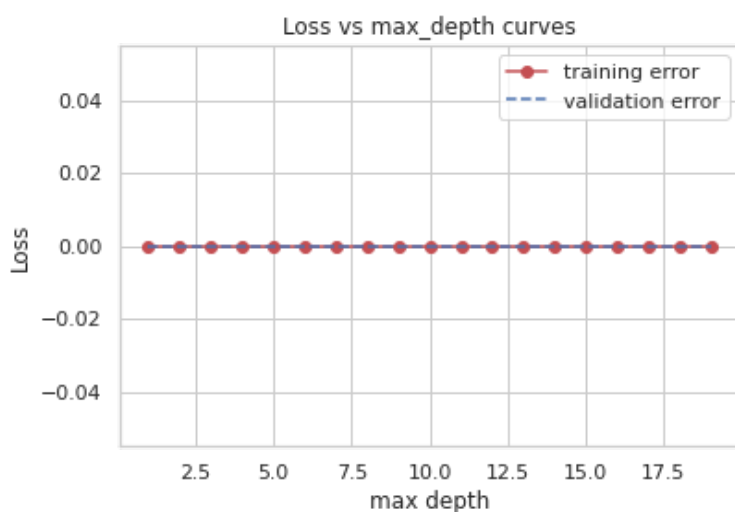
In [58]:

```
train=[]
valid=[]
for i in range(1,20):
    dec = DecisionTreeClassifier(criterion = 'gini', min_samples_leaf = 12, random_state
=3, max_depth=i)
    dec.fit(x_train, y_train)
    train.append(1- accuracy_score(dec.predict(x_train), y_train))
    valid.append(1- accuracy_score(dec.predict(x_valid), y_valid))

plt.title('Loss vs max_depth curves')
depth=[i for i in range(1,20)]
plt.plot(depth,train,'o-', color = 'r', label = 'training error')
plt.plot(depth,valid, '--', color = 'b', label = 'validation error')
plt.xlabel('max depth')
plt.ylabel('Loss')
plt.legend()
```

Out[58]:

<matplotlib.legend.Legend at 0x7f949a9c0d60>




The loss curve also shows some over-fitting, but the lowest loss is obtained for a max\_depth = 3. Now, let's apply these hyper-parameter values and let's see the model performance and generalization.

In [59]:


```
tuned_DTC = DecisionTreeClassifier(criterion='gini', max_depth = 3, min_samples_leaf = 1
2, random_state=3)
ml_model(tuned_DTC, 'Decision Tree', x_train = x_train, y_train = y_train, x_valid = x_v
alid, y_valid = y_valid, x_test = None)



generate_learning_curves(
    model = tuned_DTC,
    model_name = "Decision Trees",
    X = input_data,
    y = targets,
    ylim=(0.7, 1.01),
    epochs=5)
```

\*\*\*\*\*  
Training Set Performance:  
\*\*\*\*\*  
Accuracy of Decision Tree : 1.00  
ROC AUC Score of Decision Tree : 1.00  
Confusion Matrix of Decision Tree :



\*\*\*\*\*  
Validation Set Performance:  
\*\*\*\*\*  
Accuracy of Decision Tree : 1.00  
ROC AUC Score of Decision Tree : 1.00  
Confusion Matrix of Decision Tree :





In [60]:

```
#Section 5 : Model Performance on Test Data
```

In [61]:

```
print('For the tuned Decision Tree model :\n')
ml_model(tuned_DTC, 'tuned Decision Tree', x_train = None, x_valid = None, x_test = x_test, y_test = y_test.values.ravel() )
```

For the tuned Decision Tree model :

\*\*\*\*\*

Test Set Performance:


\*\*\*\*\*


Model Time Complexity on Test Data: 1.455 milli seconds

Accuracy of tuned Decision Tree : 1.00

ROC AUC Score of tuned Decision Tree : 1.00

Confusion Matrix of tuned Decision Tree :





In [ ]: